


COWL

Confinement with Origin Web Labels

Deian Stefan

<http://cowl.ws>

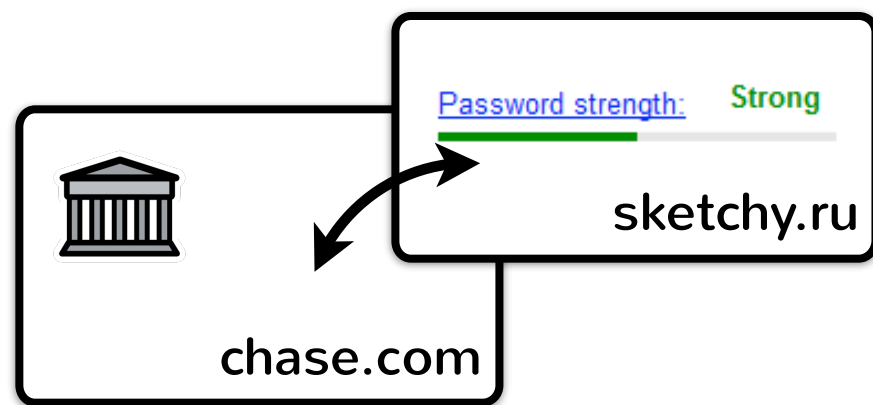
Where COWL comes from...




Crucial for the Web, but fall short in some cases...

Where does DAC fall short?


Untrusted libraries



Third-party mashups



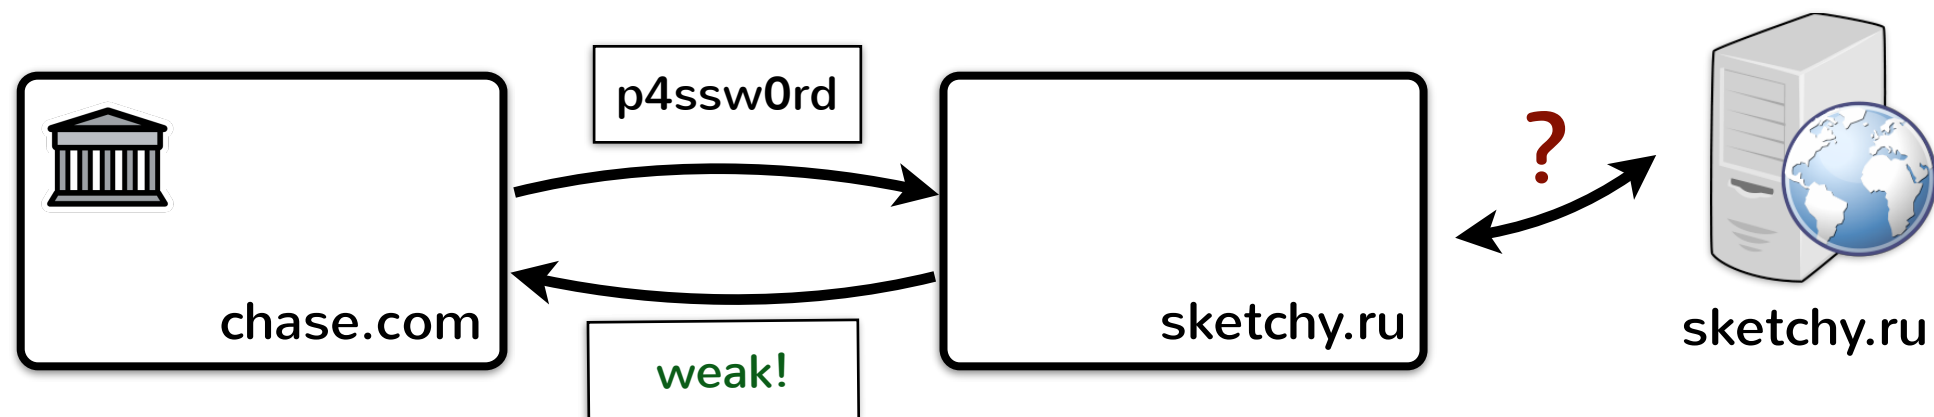
Mutually distrusting services



How does DAC fall short?

Forces choice between functionality and privacy

- E.g., password strength checker library



- **Privacy:** use CSP+sandbox to disallow communication
- **Functionality:** allow checker to fetch common pass.

How does DAC fall short?

Forces choice between functionality and privacy

- E.g., mint.com-like client-side third-party mashup



- **Privacy:** bank doesn't give mint.cc access to data
- **Functionality:** bank cedes user data to mint.cc
(or worse: user cedes bank credentials)


Why does DAC fall short?

- Fundamentally
 - Apps rely on and use third-party code
 - This code computes on sensitive data
- DAC restricts who can access data
 - Not what it can do with the data once granted access!

Confinement (at a glance)

Idea: impose restrictions on how code uses data


- E.g., it is safe to fetch list of common password before looking at password, but once password is inspected
➡ restrict communication!



Confinement (at a glance)

Idea: impose restrictions on how code uses data


- E.g., it is safe to fetch list of common password before looking at password, but once password is inspected
➡ restrict communication!



Confinement (at a glance)

Idea: impose restrictions on how code uses data

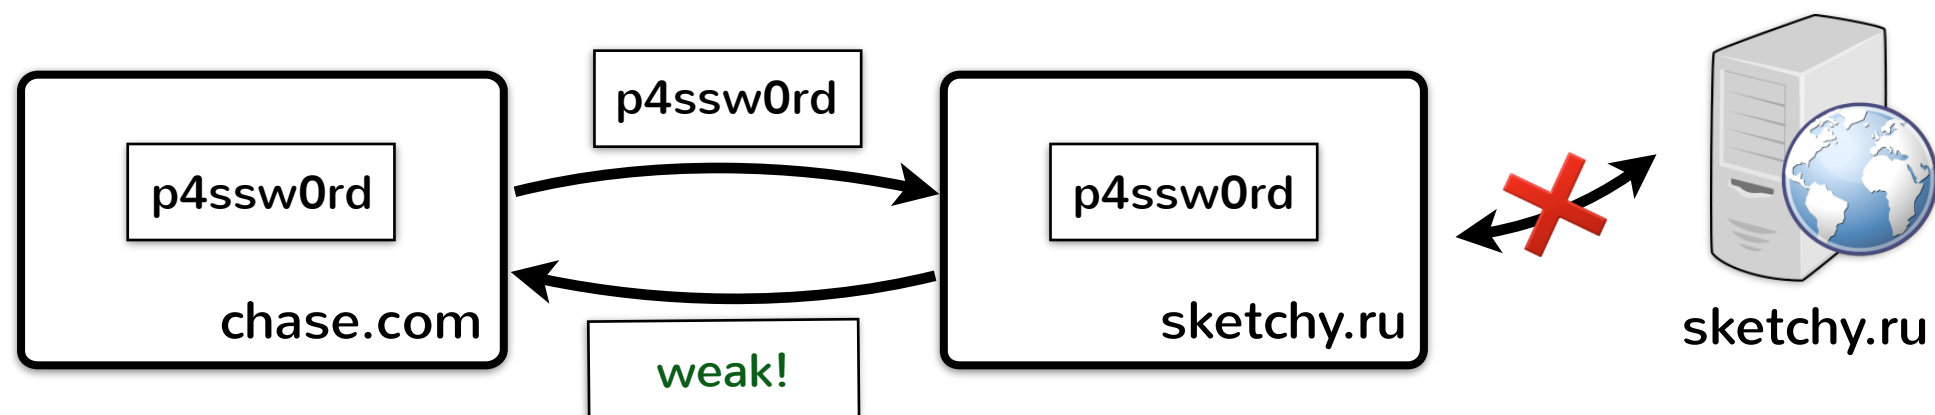
- E.g., it is safe to fetch list of common password before looking at password, but once password is inspected
➡ restrict communication!



Confinement (at a glance)

Idea: impose restrictions on how code uses data

- E.g., it is safe to fetch list of common password before looking at password, but once password is inspected
➡ restrict communication!



COWL goals

- **Allow authors to declare secrecy/integrity policy on data**
 - Ensure data is protected when shared with untrusted code (third-party or own)
- **Allow authors to consider own code untrusted**
 - I.e., allow authors to privilege separate their apps and run code with least privileges

COWL use cases

- Confining untrusted third-party services
- Sharing data with third-party mashups
- Content isolation (via privilege separation)
 - Similar to suborigin use case
- Running content with least privileges

COWL framework

- **Policy specification**
 - **Labels:** generalization of origin to conjunctions and disjunctions of origins
 - **Privileges:** makes authority of pages explicit with labels
- **Policy enforcement**
 - Enforcement of policy end-to-end via confinement

Labels

- Specifying policy in the browser
 - Labels are associated with contexts
 - Labels can be associated with clonable objects (use with postMessage and XHR)
- Specifying policy server-side
 - Labels can be associated with XHR responses

Privileges

- Explicit control over page privilege with JS
 - Page has default privilege = its origin
- Controlling default page privilege
 - E.g., to ensure page doesn't have authority of origin

Confinement enforcement

- Each context has a COWL state
 - Labels and privilege of context used to dictate who context can communicate with
- In browser: changes to WebMessaging & HTML5
- Browser-server: changes to Fetch
 - Restrict requests to prevent leaking browser data
 - Restrict responses if label of response is more sensitive than context label (non labeled-json request)

Current status

- **FPWD**

- Need to address some issues, but most features in place

- **Implementation**

- Firefox patch in the works, should be done this year
- Working on polyfill (using CSP & 3rd trusted server)
- Would love to talk to Chromium/IE teams

Issues

- What to do about same-origin, but differently labeled contexts
- Overtainting in top-level context
 - Allow leak? Leave it open-ended?
- Warnings and security reporting: what to do about bad headers

Future directions

- Light weight, in-context workers
 - Make it easier to compartmentalize apps
- Labeled-credentials?
- Clearance
 - Useful way to ensure context can never read certain kinds of data
- Covert channels
 - Enumerate leakage points (maybe as separate doc?)

Thanks!

<http://cow1.ws>